

Analysis of Neural Network Results

1. Data Preprocessing

Initial Data Setup:

Imported data: The dataset (charity_data.csv) was read into a pandas DataFrame.

Dropped non-beneficial columns: Columns like 'EIN' and 'NAME' were removed because they do not contribute to the model.

Categorical Data Processing:

APPLICATION_TYPE: I replaced application types with low frequencies (those appearing less than 500 times) with a more general category "Other". This helps reduce the dimensionality caused by infrequent categories.

CLASSIFICATION: Similarly, the "CLASSIFICATION" column was cleaned by grouping classifications that appeared less than 100 times into the "Other" category.

	APPLICATION_TYPE	AFFILIATION	CLASSIFICATION	USE_CASE	ORGANIZATION	STATUS	INCOME_AMT	SPECIAL_CONSIDERATIONS	ASK_AMT	IS_SUCCESSFUL
0	T10	Independent	C1000	ProductDev	Association	1	0	N	5000	1
1	T3	Independent	C2000	Preservation	Co-operative	1	1-9999	N	108590	1
2	T5	CompanySponsored	C3000	ProductDev	Association	1	0	N	5000	0
3	T3	CompanySponsored	C2000	Preservation	Trust	1	10000-24999	N	6692	1
4	T3	Independent	C1000	Heathcare	Trust	1	100000-499999	N	142590	1
...
34294	T4	Independent	C1000	ProductDev	Association	1	0	N	5000	0
34295	T4	CompanySponsored	C3000	ProductDev	Association	1	0	N	5000	0
34296	T3	CompanySponsored	C2000	Preservation	Association	1	0	N	5000	0
34297	T5	Independent	C3000	ProductDev	Association	1	0	N	5000	1
34298	T3	Independent	C1000	Preservation	Co-operative	1	1M-5M	N	36500179	0

34299 rows x 10 columns

Data Transformation:

One-hot encoding: Used pd.get_dummies() to convert categorical variables into numeric values, making them suitable for machine learning algorithms.

	APPLICATION_TYPE	AFFILIATION	CLASSIFICATION	USE_CASE	ORGANIZATION	STATUS	INCOME_AMT	SPECIAL_CONSIDERATIONS	ASK_AMT	IS_SUCCESSFUL
0	T10	Independent	C1000	ProductDev	Association	1	0	N	5000	1
1	T3	Independent	C2000	Preservation	Co-operative	1	1-9999	N	108590	1
2	T5	CompanySponsored	C3000	ProductDev	Association	1	0	N	5000	0
3	T3	CompanySponsored	C2000	Preservation	Trust	1	10000-24999	N	6692	1
4	T3	Independent	C1000	Heathcare	Trust	1	100000-499999	N	142590	1
...
34294	T4	Independent	C1000	ProductDev	Association	1	0	N	5000	0
34295	T4	CompanySponsored	C3000	ProductDev	Association	1	0	N	5000	0
34296	T3	CompanySponsored	C2000	Preservation	Association	1	0	N	5000	0
34297	T5	Independent	C3000	ProductDev	Association	1	0	N	5000	1
34298	T3	Independent	C1000	Preservation	Co-operative	1	1M-5M	N	36500179	0

34299 rows x 10 columns

Feature/Target Split: The target variable (IS_SUCCESSFUL) was separated from the feature variables, and the dataset was split into training and testing datasets using train_test_split.

Scaling: A StandardScaler was applied to scale the features, ensuring the neural network can learn effectively by preventing large numerical values from dominating.

2. Model Building

Model Definition:

A **Sequential neural network** was defined with:

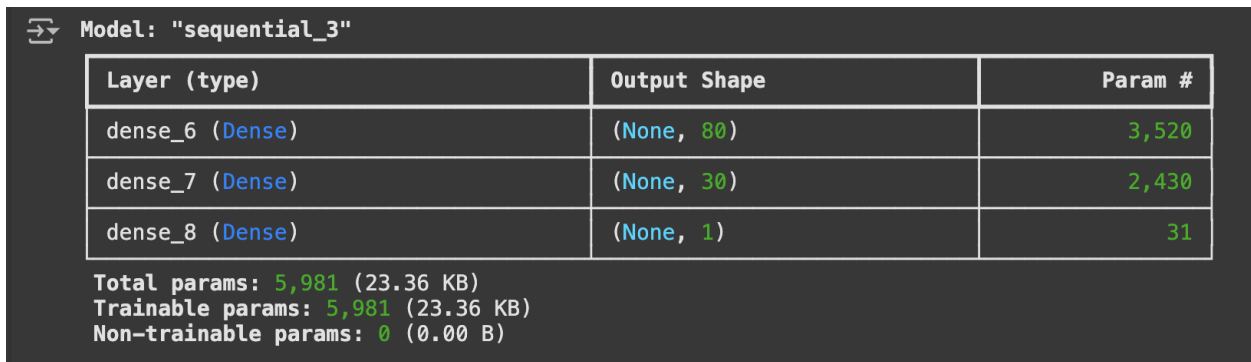
Input Layer: The number of neurons in the input layer corresponds to the number of features in the dataset.

Hidden Layers: Two hidden layers were used:

First hidden layer with 80 units and ReLU activation.

Second hidden layer with 30 units and ReLU activation.

Output Layer: A single neuron using a **sigmoid** activation function to predict the probability of success (binary classification).



The image shows a screenshot of a Keras model summary for a model named "sequential_3". It displays a table with three columns: Layer (type), Output Shape, and Param #. The layers are dense_6 (Dense) with 80 units (3,520 params), dense_7 (Dense) with 30 units (2,430 params), and dense_8 (Dense) with 1 unit (31 params). Below the table, it shows the total parameters as 5,981 (23.36 KB), trainable parameters as 5,981 (23.36 KB), and non-trainable parameters as 0 (0.00 B).

Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 80)	3,520
dense_7 (Dense)	(None, 30)	2,430
dense_8 (Dense)	(None, 1)	31

Total params: 5,981 (23.36 KB)
Trainable params: 5,981 (23.36 KB)
Non-trainable params: 0 (0.00 B)

Compilation:

The model was compiled with:

Loss function: binary_crossentropy, which is suitable for binary classification.

Optimizer: adam, a popular optimizer that adapts the learning rate during training.

Metric: accuracy, to track how well the model is performing.

3. Model Training

The model was trained for **100 epochs** with a batch size of **32** on the preprocessed training data (X_train_scaled and y_train). I also provided validation data (X_test_scaled, y_test) to monitor overfitting.

Initial Results:

During training, the accuracy improved steadily (starting from 70% and reaching close to 74% after several epochs).

The **loss** decreased while the **validation accuracy** remained around the same range (72-74%).

4. Evaluation of Results

Looking at the training results from the logs:

Accuracy fluctuated slightly, ranging from **70% to 74%** across epochs.

Validation accuracy was consistently around **72-73%** and stabilized around this level.

Loss seemed to decrease slightly, suggesting some learning took place, but the improvement was slow.

5. Analysis of Model Performance

A few points to consider about model performance:

Model underfitting: The model might be underfitting. Despite training for 100 epochs, the accuracy is relatively low (around 73%), and the model doesn't seem to improve substantially. This could indicate that the model is not complex enough to capture the patterns in the data.

Limited improvements: The validation accuracy has plateaued, and the loss remains relatively constant, which means there may be a limit to the model's performance with the current setup.

6. Next Steps for Model Improvement

Increase Model Complexity: Adding more hidden layers or increasing the number of units per layer may help the model capture more complex patterns.

Learning Rate Adjustment: Adjusting the learning rate can help the optimizer converge faster or more accurately.

More Epochs: Increasing the number of epochs might give the model more time to learn, but be careful of overfitting.

Hyperparameter Tuning: Experiment with different activation functions (like LeakyReLU), optimizers, or loss functions.

Early Stopping: Implement early stopping during training to stop the model from overfitting once the validation accuracy starts to degrade.

Cross-validation: Use k-fold cross-validation to get more robust results.

Feature Engineering: Revisit feature engineering—perhaps certain features could be combined or new features could be derived from existing ones to improve model performance.

7. Potential Data Issues

Imbalanced Classes: If the dataset is imbalanced (e.g., more successful applications than unsuccessful ones), it could skew the results. We might need to consider techniques such as class weighting, oversampling, or undersampling.

Data Leakage: Ensure that no future information (such as outcome data) is leaking into the training set.

Conclusion

The deep learning model has started to show promising results, but there is still room for optimization. Fine-tuning the model architecture and experimenting with more advanced techniques should help improve performance. Additionally, addressing potential data issues will ensure that the model generalizes well to new, unseen data.